

Validating TOSCA application topologies*

Antonio Brogi, Antonio Di Tommaso, Jacopo Soldani

Department of Computer Science, University of Pisa, Italy

Keywords: TOSCA, Validation, Cloud Application Topologies

Abstract: The *Topology and Orchestration Specification for Cloud Applications* (TOSCA) is a standardised metamodel that permits specifying cloud applications, and automating their deployment and management. TOSCA permits describing the structure of an application as a topology graph, which can then be exploited by TOSCA-compliant cloud platforms to automate the deployment of the components forming an application. In this paper we formalise the conditions that must hold for checking the validity of TOSCA application topologies, and we present an open-source prototype of validator (called SOMMELIER) that implements such conditions.

1 INTRODUCTION

Cloud computing is revolutionising IT by enabling on-demand network access to shared pools of configurable computing resources. Current cloud technologies however suffer from a lack of standardisation, with different providers offering similar resources in a different manner (Armbrust et al., 2010). Hence, how to deploy and flexibly manage complex composite applications over heterogeneous cloud platforms is one of the main challenges that have emerged after the cloud revolution (Brogi et al., 2014a).

OASIS recently released a YAML² version (OASIS, 2016) of TOSCA, the Topology and Orchestration Specification for Cloud Applications. The TOSCA Simple Profile in YAML provides a YAML-based modelling language that permits specifying portable cloud applications, and to automate their deployment and management.

TOSCA permits describing the structure of a cloud application as a typed, directed topology graph. The nodes in a topology graph represent application components (e.g., a web-based frontend, a web server, a backend database, a database management system), and each of them can be also associated with the requirements, capabilities, and management operations of the corresponding component. The arcs in a topology graph instead model relationships among application components (e.g., the frontend is deployed on the server and connected to the database, which is in turn hosted on the database management system) by

associating the requirements of a node with capabilities featured by other nodes.

TOSCA applications can be automatically deployed on TOSCA-compliant cloud platforms (Binz et al., 2013). The processing of a TOSCA application is declarative, and strictly depends on the inter-node relationships in its topology: Initially, (i) all nodes without requirements on other nodes are deployed. Then, (ii) the nodes whose requirements are actually satisfied (by capabilities offered by the nodes that have been deployed) are deployed, and their outgoing relationships are properly processed. Step (ii) is repeated until all the nodes in the application topology have been deployed (OASIS, 2013b).

Consider, for instance, the application topology in Fig. 1. The topology nodes are a frontend, a server, a database, and a database management systems, while its relationships specify that the frontend must be deployed on the server and connected to the database, and that the database must be hosted on the management system. Step (i) would result in first deploying the server and the management system. Step (ii) would then result in inserting the database among those hosted by the management system. Finally, step (ii) would be repeated to deploy the frontend on the server, and to set up the connection from the frontend to the database.




Figure 1: Example of application topology.

*This document is a preprint of the paper “Validating TOSCA application topologies” that will appear in the proceedings of MODELSWARD 2017.

²YAML Ain’t Markup Language (<http://yaml.org>).

The above (toy) example clearly illustrates how the deployment of a TOSCA application strictly depends on the inter-node relationships in its topology. This is even more evident if we consider that the processing of TOSCA nodes and relationships relies on the configuration information that they specify, and which is contextualised with constraints on how to interconnect them (OASIS, 2013b). It is hence important to ensure — at design time — that application topologies are valid, and in particular that the relationships interconnecting the nodes in such topologies have been properly specified.

In this perspective, the main contributions in this paper are twofold:

- (i) We formalise the conditions that must hold to have valid TOSCA application topologies. More precisely, since TOSCA permits specifying constraints on how to interconnect nodes (e.g., which capabilities can satisfy a requirement, or which types of relationships can be used to interconnect them), we systematically map such constraints into conditions that must hold to ensure the validity of a TOSCA application topology.
- (ii) We present SOMMELIER, which is a prototype of validator for TOSCA application topologies. SOMMELIER checks whether the topology of a TOSCA application satisfies all interconnection constraints, and, if this is not the case, it displays all violations.

Notice that SOMMELIER can be fruitfully exploited by TOSCA application developers to automatically validate the topologies of their applications at design time, as they currently have to do it manually.

It is also worth highlighting that SOMMELIER fully integrates with the OpenStack TOSCA parser (OpenStack, 2016). As both SOMMELIER and the OpenStack TOSCA parser are open-source, they can lay the foundations for an open-source toolset for supporting TOSCA application developers from the design time till the run time (Brogi et al., 2014b).

In this perspective, the output of SOMMELIER can be fruitfully exploited for improving the functionalities of TOSCA editors (e.g., the graphical editor in Alien4Cloud (Alien4Cloud, 2016)). SOMMELIER can indeed be useful for validating TOSCA application topologies while they are being developed.

The rest of this paper is organised as follows. In Sect. 2 we provide the necessary background on TOSCA, and we discuss related work. In Sect. 3 we formalise the conditions for validating TOSCA application topologies by systematically analysing the component interconnection constraints that can be ex-

pressed in TOSCA. In Sect. 4 we present SOMMELIER, and we show how it permits validating TOSCA application topologies. Finally, in Sect. 5 we draw some concluding remarks.

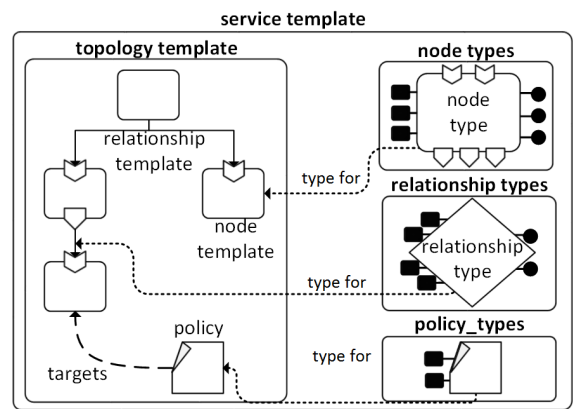
2 BACKGROUND AND RELATED WORK

2.1 Background: TOSCA

TOSCA (OASIS, 2016) is an OASIS standard whose main goals are to enable (i) the specification of portable cloud applications and (ii) the automation of their deployment and management. TOSCA provides a YAML-based, machine-readable modelling language that permits describing cloud applications. The obtained application specifications can then be processed to automate the deployment and management of the specified applications.

TOSCA permits specifying a cloud application as a `service_template`, which is in turn composed by a `topology_template`, and by the types needed to build such a `topology_template` (Fig. 2). The `topology_template` is a typed directed graph that describes the topological structure of the composite cloud application. Its nodes (called `node_templates`) model the application components, while its edges (called `relationship_templates`) model the relationships occurring among such components. The `topology_template` may also contain typed `policies`, which permit specifying non-functional information about the `node_templates` they target.

The `node_templates` and `relationship_templates` are typed by means of `node_types` and `relationship_types`, respectively. A node type de-



Legend ■Property ●Interface 🏠Capability ◇Requirement

Figure 2: TOSCA `service_template` (OASIS, 2016).

defines the observable properties and attributes of a component, its possible requirements, the capabilities it may offer to satisfy other components' requirements, and the interfaces through which it offers its management operations. Capabilities are also typed (by means of `capability_types`), to permit specifying their properties, attributes, and `valid_source_types` (viz., the node types that can be satisfied by such capabilities).

A relationship type instead describes the observable properties and attributes of a relationship occurring between two application components, and the interfaces through which it offers its management operations. A relationship type can also put constraints on its `valid_target_types` (viz., the capability types that can be targeted by a relationship whose type is that under definition).

TOSCA type system supports inheritance. A node type can be defined by extending another, thus permitting the former to inherit the latter's properties, attributes, requirements, interfaces, and operations. Analogously, a relationship type or a capability type can extend another to inherit all its features.

TOSCA also prescribes the format (called CSAR — *Cloud Service ARchive*) to archive application specifications along with the installable and executable files needed to properly instantiate the specified applications. This is because the modelling language illustrated above only allows developers to specify the application topology and its management and to give it in a *.tosca* document. Such document must be packaged together with the artefacts implementing its components so as to make all such artefacts available to the execution environment³.

2.2 Related work

Software systems are assuming a key role in everyday life, and guaranteeing that such systems satisfy their functional and non-functional requirements is becoming imperative and increasingly difficult (Marchetti, 2014). Validation techniques are hence crucial, as their purpose is precisely to evaluate software systems and to determine whether they satisfy specified requirements (Geraci, 1991).

The need for validation is also recognised by TOSCA (OASIS, 2016). Indeed, TOSCA permits specifying constraints that must be satisfied while interconnecting application components to form the topology of a cloud application. However, the current support for checking such constraints is limited,

³A more detailed, self-contained introduction to TOSCA can be found in (Brogi et al., 2014b).

which makes the process of validating TOSCA applications cumbersome and time-consuming.

The OpenStack TOSCA Parser (OpenStack, 2016) can be used to check whether TOSCA application specifications are syntactically correct. More precisely, the OpenStack TOSCA Parser permits checking that all TOSCA elements have been specified in the proper section (e.g., node types in the `node_types` section, node templates in the `node_templates` section, etc.), and that node templates and relationship templates are specified coherently with the structure given by the corresponding types (e.g., the capabilities assigned in a node template are also defined in the corresponding node type, integer properties contain integer values, etc.). However, despite the OpenStack TOSCA Parser checks that the actual value assigned to a template's field is of the proper type, it does not perform any check about the "meaningfulness" of such value. For instance, it checks whether the `node` field of a requirement assignment contains the name of a node template, but it does not check whether such node template satisfies the interconnection constraints specified in the corresponding requirement definition. Our objective is instead to permit checking such a kind of constraints.

Similar considerations apply to other TOSCA parsers currently available, e.g., the *brooklyn-tosca* parser (CloudSoft, 2016), or the TOSCA parser in Alien4Cloud (Alien4Cloud, 2016).

TOSCA is also available in an XML-based version (OASIS, 2013a), which still permits specifying the topology of a composite cloud application, as well as the constraints to interconnect the components forming such topology. TOSCA XML also permits specifying management plans, in the form of workflows orchestrating the management operations offered by the application components.

TOSCA XML applications can be edited with the Winery graphical environment (Kopp et al., 2013), and executed with the OpenTOSCA engine (Binz et al., 2013). However, despite Winery and OpenTOSCA are equipped with parsers checking whether the XML sources of applications are syntactically correct, there is no support for validating the interconnections forming application topologies.

(Hirmer et al., 2014) propose an approach for automatically completing TOSCA XML application topologies. The approach is based on some of the interconnection constraints that can be expressed in TOSCA XML, which are exploited to decide which available components can satisfy a dangling requirement of a component. The approach of (Hirmer et al., 2014) is a first approach exploiting the interconnection constraints that can be expressed in TOSCA

XML. The aim of (Hirmer et al., 2014) is however different from ours, as they rely on TOSCA XML, and since they consider only some of the interconnection constraints that can be expressed in TOSCA XML (viz., those constraining the relationship types and capability types that can be validly used to fulfil a requirement). Our aim is instead to systematically map all constraints that can be expressed in TOSCA to formal conditions, hence enabling a full validation of TOSCA application topologies.

On the other hand, it is possible to validate management plans in TOSCA XML. Management protocols (Brogi et al., 2015; Brogi et al., 2016) permit specifying the management behaviour of the components forming an application, which can then be combined (according to the application topology) to obtain the management behaviour of the whole application. Such a behaviour permits automating various useful analyses (e.g., determining the validity and effects of a management plan, or which plans reach a given goal). However, despite topologies are crucial to determine the management behaviour of applications, they are assumed to be valid, because of the lack of support for validating application topologies.

In summary, despite TOSCA is designed to permit specifying constraints on how to interconnect application components to form application topologies, the support for checking such constraints is currently lacking. Existing parsers can be used to check whether TOSCA applications are syntactically correct, but currently there are no tools that permit validating TOSCA application topologies. Our objective is to address this lack, by systematically mapping the interconnection constraints that can be expressed in TOSCA to formal conditions, and by exploiting such conditions to devise a first support for validating TOSCA application topologies.

Finally, it is worth mentioning that our approach shares its baselines with most of existing approaches for checking the validity of component-based systems at design time, e.g., (Speck et al., 2002; Caporuscio et al., 2004; Wu et al., 2003), or for automatically obtaining valid component-based systems, e.g., (Autili et al., 2007; Pelliccione et al., 2008). Given a specification of a component-based system (a TOSCA application specification, in our case), we try to enforce that all interconnections between components are properly settled (which in our case means to check whether the topology of an application satisfies all the specified interconnection constraints). The main difference between such approaches and ours is given by the context: (Speck et al., 2002; Caporuscio et al., 2004; Wu et al., 2003; Autili et al., 2007; Pelliccione et al., 2008) focus on systems that have to offer a cer-

tain functionality by validly composing the functionalities offered by their components. We instead focus on ensuring that the dependencies between the components forming a TOSCA application are properly specified, as such information is the basis for orchestrating the management of a TOSCA application.

3 VALIDATING TOPOLOGIES

The objective of this paper is to permit validating the topology of a TOSCA application (viz., its `topology_template`). As shown in Sect. 2.1, the `topology_template` is a typed directed graph whose nodes model the application components, and whose edges model the relationships occurring between such components. Each relationship specifies that a requirement of the source node must be actually satisfied by (a capability of) the target node.

As we anticipated in Sect. 1, in order to validate the topology of a TOSCA application, we must check whether all inter-component relationships have been properly specified. For each relationship, we must validate all the TOSCA elements concerning its source, the relationship itself, and its target. In this respect, please note that the source of a relationship is always a requirement of a node, that a relationship is actually a relationship template, and that the target of a relationship is either a node or a capability of a node (OASIS, 2016).

We hereafter *systematically map*⁴ the interconnection constraints expressed while defining TOSCA requirements, relationships, and capabilities to formal conditions. We separately discuss the conditions to validate the interconnection constraints specified in the sources of relationships (Sect. 3.1), in relationships (Sect. 3.2), and in the targets of relationships (Sect. 3.3). We then combine all such conditions to define the notion of validity for a TOSCA application topology (Sect. 3.4).

3.1 Validating sources of relationships

We hereby illustrate how to validate the source of each relationship appearing in an application topology. As the source of a relationship is always a node's requirement (OASIS, 2016), we must check that each

⁴After systematically reading the whole TOSCA specification (OASIS, 2016), we have extracted the portions that permit specifying interconnection constraints. We hereafter recall such portions, and we show how to directly map them to formal conditions that must hold to have valid TOSCA application topologies.

```
req_name: cap_type_name
```

(a)

```
req_name:
  capability: cap_type_name
  node: node_type_name
  relationship: rel_type_name
  occurrences: [ min_occ, max_occ ]
```

(b)

Figure 3: (a) Simple and (b) extended grammars for *requirement definitions* (OASIS, 2016).

relationship outgoing from a requirement is compatible with the requirement definition itself.

In the following, we recall how to declare a requirement in TOSCA, and what is its meaning according to the TOSCA specification. We then single out the checks that must be performed to ensure that no relationship is violating the constraints given by a requirement declaration.

3.1.1 How to *define* a requirement in TOSCA

We hereby recall how to define a requirement in a TOSCA node type, and how to instantiate a requirement in a TOSCA node template.

A node type n_{type} defines the set of named requirements that can be exposed by a node template of such type. Each *requirement definition* can be given within the `requirements` section of n_{type} according to the (a) simple or (b) extended grammars in Figure 3. In both cases, a developer can define the requirement name (`req_name`), and she can provide constraints for its satisfaction.

The (a) simple grammar requires to indicate the name (`cap_type_name`) of a valid capability type that can fulfil the requirement under definition. The (b) extended grammar also permit indicating three optional constraints.

- `node` permits indicating the name (`node_type_name`) of a valid node type that contains a capability definition that can be used to fulfil the requirement under definition.
- `relationship` permits providing the name of a relationship type that can be validly used to build a relationship template when fulfilling the requirement.
- `occurrences` permits instructing the minimum (`min_occ`) and maximum (`max_occ`) occurrences of the requirement under definition in a node template of the corresponding node type⁵.

⁵Since the focus of this paper is on validating inter-

```
req_name: node_temp_name
```

(a)

```
req_name:
  node: node_temp_name | node_type_name
  relationship: rel_temp_name | rel_type_name
  capability: cap_name | cap_type_name
  node_filter: node_filter_definition
```

(b)

Figure 4: (a) Simple and (b) extended grammars for *requirement assignments* (OASIS, 2016).

By default, `min_occ` and `max_occ` are both set to 1.

A node template n_{temp} instantiates a node type n_{type} in a TOSCA application topology, and it permits declaring the requirements actually exposed by the application component corresponding to such node template. Such requirements can be given through so-called *requirement assignments* within the `requirements` section of n_{temp} . Each requirement assignment instantiates a corresponding⁶ requirement definition in n_{type} . The (a) simple and (b) extended grammars for requirement assignments are displayed in Figure 4.

The (a) simple grammar permits indicating only the name (`node_temp_name`) of the concrete node template that is actually satisfying the requirement under assignment. This notation is only valid if the corresponding requirement definition (in the node type of the node template that is being specified) indicates at least a valid capability type that can be found in the target node template.

The (b) extended grammar can be used if the requirement assignment has to provide more information than just the name of the target node template.

- `node` is used to identify the target node of a relationship. It can be used to provide either the name of the node template that is actually fulfilling the requirement under assignment, or the name of a node type that a TOSCA-compliant cloud provider will use to select a type-compatible node template to fulfil the requirement at run-time.
- `relationship` is an optional reserved keyword that can be used to provide either the name of a

component dependencies in TOSCA application topologies, and since `occurrences` is not giving any constraint concerning inter-component dependencies, we shall not formalise the trivial condition to validate `occurrences`.

⁶A node template's requirement assignment corresponds to a node type's requirement definitions if they have the same name `req_name`.

relationship template (to relate the source node to the — capability in the — target node when fulfilling the requirement), or the name of a relationship type (that will be used by a TOSCA-compliant cloud provider to select a type-compatible relationship template to relate the source and target nodes at run-time).

- `node_filter` permits defining additional constraints that TOSCA-compliant cloud providers will use to select a type-compatible target node that can fulfil the requirement under assignment at run-time.

3.1.2 How to *validate* a TOSCA requirement

We hereby single out the conditions that must hold to validate sources of relationships at design-time⁷. In doing so, we exploit the following shorthand notation.

Notation. *In the following, we write:*

- $\text{name}(\cdot)$ and $\text{type}(\cdot)$ to denote the name and type associated to a TOSCA element, respectively,
- $t' \geq t$ to denote that t' extends⁸ or is equal to t ,
- $e.\varepsilon$ to denote the field ε of the TOSCA element e (which is \perp if ε is not defined in e), and
- $R(\cdot)$ and $C(\cdot)$ to denote the requirements and capabilities assigned by a node template or defined in a node type.

Given a TOSCA application topology, the sources of its inter-component relationships are valid if each of its node templates satisfies all the aforementioned conditions.

Condition 1. *Let n_{temp} be a node template of type*

⁷We hence exclude all constraints that permit defining how to automatically complete the topology of a TOSCA application (by selecting type-compatible node templates and relationship templates to satisfy pending requirements — e.g., `node: node_type_name, relationship: rel_type_name`, and `capability: cap_type_name` in Figure 4.(b)). Our validation approach can however be useful for double-checking automatically completed TOSCA application topologies (as well as for helping the automatic completion itself).

⁸Given that t and t' are TOSCA types, t' extends t if t' is (directly or indirectly) derived from t .

n_{type} . Then:

$\forall r_a \in R(n_{temp}), \exists r_d \in R(n_{type})$:

- 1) $\text{name}(r_a) = \text{name}(r_d) \wedge$
- 2) $r_d.\text{node} \neq \perp \Rightarrow \text{type}(r_a.\text{node}) \geq r_d.\text{node} \wedge$
- 3) $r_a.\text{capability} \neq \perp \Rightarrow$
 $\text{type}(r_a.\text{capability}) \geq r_d.\text{capability} \wedge$
- 4) $r_a.\text{capability} = \perp \Rightarrow$
 $\exists c \in C(r_d.\text{node}) : \text{type}(c) \geq r_d.\text{capability} \wedge$
- 5) $r_d.\text{relationship} \neq \perp \wedge r_a.\text{relationship} \neq \perp \Rightarrow$
 $\text{type}(r_a.\text{relationship}) \geq r_d.\text{relationship}$

Firstly, we must check that, for each requirement assigned r_a in a node template n_{temp} (of type n_{type}), there exists a corresponding requirement definition r_d in n_{type} (Condition 1.1).

We can then check that no requirement assignment r_a is violating the constraints indicated by the corresponding requirement definition r_d :

- The (optional) field `node` of a requirement definition r_d permits indicating the name of a valid node type that can be used to fulfil a requirement. Hence, if `node` is specified in r_d , the type of the node template targeted by each corresponding requirement assignment r_a has to extend or to be equal to that indicated in `node` (Condition 1.2.)
- A requirement definition r_d indicates (either inline or with the keyword `capability`) the name of a valid capability type that can fulfil the requirement. This is ensured if each corresponding requirement assignment r_a directly targets a capability whose type extends or is equal to that indicated in the requirement definition (Condition 1.3). Alternatively, if a requirement assignment r_a is only indicating the target node template (but not pointing to any of its capabilities), we must check whether such node template is offering at least one type-compatible capability (Condition 1.4).
- Finally, the (optional) field `relationship` of a requirement definition r_d permits indicating the name of a relationship type that can be validly used to build a relationship template when fulfilling the requirement. Hence, if `relationship` is specified in r_d , the type of a relationship template outgoing from a corresponding requirement assignment r_a (if any) extends or is equal to that indicated in `relationship` (Condition 1.5).

3.2 Validating relationships

We hereby illustrate how to validate the relationships occurring between the nodes appearing in an applica-

tion topology. As such relationships are expressed by typed relationship templates (OASIS, 2016), we must check that each relationship template is instantiated without violating any of the constraints given in the corresponding relationship type.

In the following, we recall how to declare a relationship type in TOSCA, and which is its meaning according to the TOSCA specification. We then illustrate the conditions that must hold to ensure that no relationship template is violating the constraints given by the corresponding relationship type.

3.2.1 How to *define* a relationship in TOSCA

A relationship template instantiate a relationship type to specify the actual occurrence of a manageable relationship between two node templates in an application topology. A relationship type can be defined with the grammar in Figure 5, which permits specifying the name (*rel_type_name*) of a relationship type, as well as the features and constraints characterising such relationship type.

```

rel_type_name:
  derived_from: parent_rel_type_name
  version: version_number
  description: rel_description
  properties: property_definitions
  attributes: attribute_definitions
  interfaces: interface_definitions
  valid_target_types: [ cap_type_names ]

```

Figure 5: Grammar for *relationship types* (OASIS, 2016).

- The (optional) field *derived_from* permits indicating the name (*parent_rel_type_name*) of the parent relationship type⁹. The relationship type under definition inherits all the features the parent relationship type, and it can override some of them (OASIS, 2013b).

For instance, if the relationship type under definition does not specify a new list of *valid_target_types*, then it takes that of the parent relationship type. Otherwise, the parent's list of *valid_target_types* is overridden by that specified in the relationship type under definition.

- *version* and *description* are optional fields that permit versioning and describing (in natural language) a relationship type.
- The (optional) fields *properties* and *attributes* permit describing the structure

⁹All TOSCA relationship types should be derived (directly or indirectly) from the `tosca.relationships.Root` relationship type (OASIS, 2016).

of the desired and actual state of a relationship, respectively. Concrete state values will be then given in the relationship templates instantiating the relationship type under definition.

- *interfaces* is a (optional) field that permits declaring the management operations that can be offered by relationship templates whose relationship type is that under definition.
- *valid_target_types* is a (optional) list of one or more names (*cap_type_names*) of capability types that are valid targets for the relationship under definition.

3.2.2 How to *validate* a TOSCA relationship

Since a relationship template instantiates a relationship type in an application topology, it has to not violate the constraints given when defining the corresponding relationship type.

Notice that the only constraints that can be given while defining a relationship type are those concerning its *valid_target_types*. The latter can either be specified inline or inherited from the parent type.

Notation. In the following, we write $T(\cdot)$ to denote the set of capability types that are valid targets for a relationship type. Given a relationship type rel_{type} :

- If $rel_{type}.valid_target_types \neq \perp$, then $T(rel_{type})$ is the set containing all elements in $rel_{type}.valid_target_types$,
- otherwise, if $rel_{type}.derived_from \neq \perp$, then $T(rel_{type}) = T(rel_{type}.derived_from)$,
- otherwise, $T(rel_{type})$ is the set containing all capability types (meaning that all capability types are valid targets for rel_{type}).

Given a TOSCA application topology, its relationship templates are valid if they satisfy the following condition.

Condition 2. Let rel_{temp} be a relationship template of type rel_{type} . If there exists a requirement assignment r_a that is source of rel_{temp} , then:

- 1) $r_a.capability \neq \perp \Rightarrow$
 $\exists c_{type} \in T(rel_{type}) : type(r_a.capability) \geq c_{type} \wedge$
- 2) $r_a.capability = \perp \Rightarrow$
 $\exists c_a \in C(r_a.node) :$
 $\exists c_{type} \in T(rel_{type}) : type(c_a) \geq c_{type}$

Namely, if a relationship template rel_{temp} is targeting a specific capability¹⁰, then the type of such capa-

¹⁰Please recall that a relationship is outgoing from a requirement assignment r_a . The latter can either indicate the specific capability satisfying r_a , or a node template offering (at least) a capability satisfying r_a (see Figure 4).

bility has to be compatible with at least one of those indicated in `valid_target_types` (Condition 2.1).

If the source of a relationship template rel_{temp} is instead only indicating the target node template (but not pointing to any of its capabilities), we must check whether such node template is offering at least one type-compatible capability (Condition 2.2).

3.3 Validating targets of relationships

While the source of a relationship is always a requirement, the target of a relationship can be either a capability or a node template. In the latter case, the node template must offer at least a capability capable of satisfying the source requirement (OASIS, 2016).

In the following we recall how to specify a capability in TOSCA, and which is its meaning according to the TOSCA specification. We then illustrate the conditions that must hold to ensure that no relationship is violating the constraints given by its target.

3.3.1 How to define a capability in TOSCA

A capability type is a reusable entity that describes a kind of capability that a node type can declare to expose (OASIS, 2016). A capability type can be defined according to the grammar in Figure 6, which permits specifying the name (*cap_type_name*) of a capability type, as well as with the features and constraints characterising such capability type.

```
cap_type_name:
  derived_from: parent_cap_type_name
  version: version_number
  description: capability_description
  properties: property_definitions
  attributes: attribute_definitions
  valid_source_types: [ node_type_names ]
```

Figure 6: Grammar for *capability types* (OASIS, 2016).

- The (optional) field `derived_from` permits indicating the name (*parent_cap_type_name*) of the parent capability type¹¹. The capability type under definition inherits all the features the parent capability type, and it can override some of them (OASIS, 2013b).

For instance, if the capability type under definition does not specify a new list of `valid_source_types`, then it takes that of the parent capability type. Otherwise, the parent's list of `valid_so-`

¹¹ All TOSCA capability types should be derived (directly or indirectly) from the `tosca.capability.Root` capability type (OASIS, 2016).

```
cap_name: cap_type_name
```

(a)

```
cap_name:
  type: cap_type_name
  description: capability_description
  properties: property_definitions
  attributes: attribute_definitions
  valid_source_types: [ node_type_names ]
```

(b)

Figure 7: (a) Simple and (b) extended grammars for *capability definitions* (OASIS, 2016).

`urce_types` is overridden by that specified in the capability type under definition.

- `version` and `description` are optional fields that permit versioning and describing (in natural language) a capability type.
- The (optional) fields `properties` and `attributes` permit describing the structure of the desired and actual state of a capability, respectively. Concrete state values will be then given by the node templates instantiating capabilities whose type is that under definition.
- `valid_source_types` is a (optional) list of one or more names (*node_type_names*) of node types that are valid sources for relationships whose target capability is of the type under definition.

Capability types are then referred by node types to define the set of named capabilities that can be exposed by a node template of such type. Each *capability definition* can be given within the *capabilities* section of a node type n_{type} according to the (a) simple or (b) extended grammars in Figure 7. In both cases, a developer can define the capability name (*cap_name*), and she can specify the type (*cap_type_name*) of the capability under definition. With the extended grammar, a developer can also specify the (optional) fields `description`, `properties`, `attributes`, and `valid_source_types` (whose meaning is analogous to that of their homonym fields in the grammar for specifying capability types — see Figure 6).

A node template n_{temp} instantiates a node type n_{type} in a TOSCA application topology, as well as all the capabilities it defines. A node template can also assign concrete values to the properties and attributes of such capabilities, to provide additional information concerning their desired and actual state, respectively. Such assignments can be given through so-called *capability assignments* within the *capabilities* section of n_{temp} (Figure 8).

Validating sources of relationships (Condition 1)	Let n_{temp} be a node template of type n_{type} . Then, $\forall r_a \in R(n_{temp}), \exists r_d \in R(n_{type})$: (1.1) $name(r_a) = name(r_d) \wedge$ (1.2) $r_d.node \neq \perp \Rightarrow type(r_a.node) \geq r_d.node \wedge$ (1.3) $r_a.capability \neq \perp \Rightarrow type(r_a.capability) \geq r_d.capability \wedge$ (1.4) $r_a.capability = \perp \Rightarrow \exists c \in C(r_a.node): type(c) \geq r_d.capability \wedge$ (1.5) $r_d.relationship \neq \perp \wedge r_a.relationship \neq \perp \Rightarrow type(r_a.relationship) \geq r_d.relationship$
Validating relationships (Condition 2)	Let rel_{temp} be a relationship template, and let rel_{type} be its relationship type. If there exists a requirement assignment r_a that is source of rel_{temp} , then: (2.1) $r_a.capability \neq \perp \Rightarrow \exists c_{type} \in T(rel_{type}): type(r_a.capability) \geq c_{type} \wedge$ (2.2) $r_a.capability = \perp \Rightarrow \exists c_a \in C(r_a.node), c_{type} \in T(rel_{type}): type(c_a) \geq c_{type}$
Validating targets of relationships (Condition 3)	Let c_a be a capability assignment, whose corresponding type and definition are c_{type} and c_d . For each node template n_{temp} having a requirement assignment r_a such that $r_a.capability = c_a$: (3.1) $\exists n_{type} \in S(c_{type}): type(n_{temp}) \geq n_{type} \wedge$ (3.2) $\exists n_{type} \in S(c_d): type(n_{temp}) \geq n_{type}$

Table 1: Conditions to validate inter-component relationships in TOSCA application topologies.

cap.name: properties: <i>property_assignments</i> attributes: <i>attribute_assignments</i>

Figure 8: Grammar for *capability assignments* (OASIS, 2016).

3.3.2 How to validate a TOSCA capability

While defining a capability in TOSCA, the only constraints that can be given are those concerning the *valid_source_types*. Such constraints can be specified within the specification of a capability type, or within the capability definitions of a node type. The constraints given while defining a capability type can either be specified inline or they can be inherited from the parent capability type.

Notation. In the following, we write $S(\cdot)$ to denote the set of node types that are valid sources for a capability type. Given a capability type c_{type} :

- If $c_{type}.valid_source_types \neq \perp$, then $S(c_{type})$ is the set containing all elements in $c_{type}.valid_source_types$,
- otherwise, if $c_{type}.derived_from \neq \perp$, then $S(c_{type}) = S(c_{type}.derived_from)$,
- otherwise, $S(c_{type})$ is the set containing all node types (meaning that all node types are valid targets for c_{type}).

Given a capability assignment c_a , in a TOSCA application topology there might be multiple node templates (having requirement assignments) that are sources of relationship templates targeting c_a . All

such node templates must not violate any of the constraints imposed by c_a :

- Neither those indicated by its capability type c_{type} (Condition 3.1),
- nor those specified within the corresponding capability definition c_d (Condition 3.2).

Condition 3. Let c_a be a capability assignment, whose corresponding type and definition are c_{type} and c_d . For each node template n_{temp} having a requirement assignment r_a such that $r_a.capability = c_a$:

- 1) $\exists n_{type} \in S(c_{type}): type(n_{temp}) \geq n_{type} \wedge$
- 2) $\exists n_{type} \in S(c_d): type(n_{temp}) \geq n_{type}$

3.4 Valid TOSCA application topologies

In Sects. 3.1, 3.2, and 3.3 we have illustrated the conditions that permit validating sources, instances, and targets of TOSCA relationships, respectively. The following definition gathers all such conditions to permit validating all inter-component relationships appearing in a TOSCA application topology.

Definition 1. A TOSCA application topology is valid if all its node templates and relationship templates satisfy Conditions 1, 2, and 3.

For the convenience of readers, Conditions 1, 2, and 3 are recapped in Table 1.




Figure 9: Open-source toolchain for validating TOSCA application topologies.

4 THE SOMMELIER PROTOTYPE

We hereby illustrate SOMMELIER, an open-source¹² prototype for validating TOSCA application topologies (as we discussed in the previous section). The prototype of SOMMELIER is implemented in Python, and it is fully integrated with the OpenStack TOSCA parser (OpenStack, 2016).

Figure 9 illustrates how SOMMELIER performs the validation of a TOSCA application topology:

- ❶ SOMMELIER inputs a CSAR package or a .tosca document containing the TOSCA application topology to be validated.
As the current prototype of SOMMELIER is a command-line tool, the CSAR or .tosca files can be passed to SOMMELIER with a dedicated command line option.
- ❷ The input is forwarded by SOMMELIER to the *OpenStack Parser*.
- ❸ The *OpenStack Parser* checks whether the TOSCA application is syntactically correct, i.e. it checks whether all TOSCA elements have been specified in the proper section (e.g., node types in the `node_types` section) and that each template respects the structure given by the corresponding type (e.g., the capabilities assigned in a node template are also defined in the corresponding node type). If this is the case, the *OpenStack Parser* generates a representation of the TOSCA application in Python, according to its object-model (OpenStack, 2016).
- ❹ The Python objects representing the TOSCA application are returned to SOMMELIER.

¹²The source code of SOMMELIER is publicly available on GitHub at <https://github.com/di-unipi-socc/Sommelier>.




Figure 10: Structure of the Python dictionary containing the results of the validation.

- ❺ SOMMELIER validates the topology of the TOSCA application by checking whether all (the Python objects representing) its nodes and relationships satisfy all conditions¹³ in Table 1.

SOMMELIER then generates a Python dictionary¹⁴ whose abstract structure is illustrated in Figure 10. Namely, SOMMELIER associates each requirement of each node template with a list containing all validation errors affecting the relationship outgoing from such requirement. Each error is in turn represented as a list, whose first element is a numeric code (e.g., 1.2 if the requirement is violating Condition 1.2), and whose remaining elements may be used to provide further information about the error (e.g., the type and name of the target node, which is not satisfying Condition 1.2).

- ❻ The results of the validation are displayed by SOMMELIER. The current prototype of SOMMELIER either states that the service template's topology is valid, or it lists all the detected validation errors (by also explaining why such errors occurred).

Example. We now illustrate how to use SOMMELIER to validate existing TOSCA application topologies. As SOMMELIER is a command-line tool, running an instance of SOMMELIER just requires to type the following command in a command shell:

```
$ python sommelier.py
   --template-file=template-file-path
```

¹³SOMMELIER is distributed along with unit tests illustrating that it is capable of recognising the violation of all such conditions. The unit tests are publicly available within the `tests` folder in the SOMMELIER Github repository (i.e., at <https://github.com/di-unipi-socc/Sommelier/tree/master/tests>).

¹⁴<https://docs.python.org/3/tutorial/datastructures.html>.

```
$ python sommelier.py --template-file=
../tosca-parser/toscaparser/tests/data
/topology_template/tosca_elk.yaml
```

```
The application topology is valid.
```

(a)

```
$ python sommelier.py --template-file=
../tosca-parser/toscaparser/tests/data
/topology_template/transactionssubsys
tems.yaml
```

```
NODE_TEMPLATE: app
REQUIREMENT: host
1.2 - NODE_TYPE_NOT_COHERENT: The type
"tosca.nodes.WebServer" of the target
node "websrv" is not valid (as it diff
ers from that indicated in the require
ment definition).
```

(b)

Figure 11: Example of runs of SOMMELIER.

(where `template-file-path` is the absolute path where the TOSCA file to be validated is stored).

Concrete examples of run of SOMMELIER are shown in Figure 11. In the figure, the application specification that are passed to SOMMELIER are taken among those available in the GitHub repository of the OpenStack TOSCA parser. Figure 11.(a) shows that the TOSCA application topology specified in *tosca_elk.yaml*¹⁵ is valid.

Figure 11.(b) instead shows that the TOSCA application topology specified in *transactionssubsystem.yaml*¹⁶ is not valid, because the relationship outgoing from requirement *host* of the node template *app* is violating Condition 1.2 (since it targets the node template *websrv*, whose type is not compatible with those indicated in the requirement definition corresponding to *host*). Such a kind of inconsistencies would be really hard to be manually detected, without the availability of a tool like SOMMELIER.

¹⁵The *tosca_elk.yaml* file has been developed within the OpenStack TOSCA parser community, and it is publicly available at https://raw.githubusercontent.com/openstack/tosca-parser/master/toscaparser/tests/data/tosca_elk.yaml (last accessed on November 13th, 2016).

¹⁶The *transactionssubsystem.yaml* file has been developed within the OpenStack TOSCA parser community, and it is publicly available at https://raw.githubusercontent.com/openstack/tosca-parser/master/toscaparser/tests/data/topology_template/transactionssubsystem.yaml (last accessed on November 13th, 2016)

5 CONCLUDING REMARKS

TOSCA provides a standardised language that permits specifying portable cloud applications, and to automate their deployment and management. As the automated deployment of TOSCA applications is based on their topologies (OASIS, 2013b), it is fundamental to ensure that such topologies are valid.

TOSCA permits specifying the constraints that must be satisfied while interconnecting application components to form the topology of a cloud application. In this paper:

- (i) We have systematically mapped such constraints to formal conditions that must hold to ensure the validity of a TOSCA application topology. The obtained conditions are recapped in Table 1.
- (ii) We have also presented SOMMELIER, an open-source prototype of validator for TOSCA application topologies.

SOMMELIER implements all the aforementioned validation conditions (as one can readily check by running the unit tests available in its repository). SOMMELIER can hence be fruitfully exploited by TOSCA application developers to automatically validate the topologies of their applications.

SOMMELIER is fully integrated with the OpenStack TOSCA parser (OpenStack, 2016). As they are both open-source, they can lay the foundations for an open-source toolset for supporting TOSCA application developers. This is in line with the research directions indicated in (Brogi et al., 2014b), and in particular with the development of tools to support TOSCA developers from the design time till the run time.

It is worth noting that SOMMELIER can be fruitfully exploited for improving the functionalities of existing graphical editors (such as that in the Alien4Cloud platform (Alien4Cloud, 2016), for instance). SOMMELIER can indeed be useful for validating TOSCA application topologies while they are being developed. Its output could also be exploited for suggesting how to fix errors (e.g., which capability can actually satisfy a misplaced requirement) or to drive the development itself (e.g., warning the user if she is connecting a requirement to a wrong capability — or impeding her to do so). As part of our future work, we plan to integrate SOMMELIER with a graphical editor, with the long-term objective of providing a full-fledged support for the development of TOSCA applications (Brogi et al., 2014b).

It is also worth highlighting that the elements that can be expressed in TOSCA YAML (OASIS, 2016) are a subset of those that can be defined in TOSCA XML (OASIS, 2013a), as the former is a simplified

profile of the latter. We plan to investigate whether the validation conditions contained in this paper are enough to validate also the topologies of TOSCA XML applications, and to (extend and) adapt them to work with TOSCA XML, if necessary. This would permit providing the OpenTOSCA open-source environment (Kopp et al., 2013; Binz et al., 2013) also with validation capabilities. It would also permit refining the automated topology completion approach by Hirmer et al. (Hirmer et al., 2014), by ensuring the validity of obtained topologies.

Finally, as TOSCA permits describing not only the technological requirements of application components, but also their non-functional requirements (e.g., scalability constraints, in terms of minimum and maximum amount of occurrences of a node replica, or QoS requirements specified in policies applied to nodes), TOSCA applications should also be validated also from a non-functional perspective. We plan to systematically map the non-functional constraints of that can be expressed in TOSCA to formal conditions that will permit carrying out such a validation.

ACKNOWLEDGEMENTS

Work partly supported by the project *Through the fog* (PRA_2016_64) funded by the University of Pisa.

REFERENCES

- Alien4Cloud (2016). Alien4cloud. <https://github.com/alien4cloud/alien4cloud>.
- Armbrust, M., Fox, A., Griffith, R., Joseph, A. D., Katz, R., Konwinski, A., Lee, G., Patterson, D., Rabkin, A., Stoica, I., and Zaharia, M. (2010). A view of cloud computing. *Commun. ACM*, 53(4):50–58.
- Autili, M., Inverardi, P., Navarra, A., and Tivoli, M. (2007). Synthesis: A tool for automatically assembling correct and distributed component-based systems. In *Proc. of the 29th International Conference on Software Engineering*, ICSE '07, pages 784–787. IEEE Computer Society.
- Binz, T., Breitenbücher, U., Haupt, F., Kopp, O., Leymann, F., Nowak, A., and Wagner, S. (2013). Open-tosca – a runtime for toasca-based cloud applications. In *Proc. of the 11th International Conference on Service-Oriented Computing*, ICSOC'13, pages 692–695. Springer.
- Brogi, A., Canciani, A., and Soldani, J. (2015). Modelling and analysing cloud application management. In Dustdar, S., Leymann, F., and Villari, M., editors, *Service Oriented and Cloud Computing: 4th European Conference, ESOC 2015, Taormina, Italy, September 15-17, 2015, Proceedings*, pages 19–33. Springer International Publishing.
- Brogi, A., Canciani, A., Soldani, J., and Wang, P. (2016). A Petri net-based approach to model and analyze the management of cloud applications. In Koutny, M., Desel, J., and Kleijn, J., editors, *Transactions on Petri Nets and Other Models of Concurrency XI*, LNCS Transactions on Petri Nets and Other Models of Concurrency, pages 28–48. Springer Berlin Heidelberg.
- Brogi, A., Carrasco, J., Cubo, J., D'Andria, F., Ibrahim, A., Pimentel, E., and Soldani, J. (2014a). EU Project SeaClouds - Adaptive management of service-based applications across multiple clouds. In *CLOSER 2014 - Proceedings of the 4th International Conference on Cloud Computing and Services Science*, pages 758–763. SciTePress.
- Brogi, A., Soldani, J., and Wang, P. (2014b). TOSCA in a nutshell: Promises and perspectives. In Villari, M., Zimmermann, W., and Lau, K.-K., editors, *Service-Oriented and Cloud Computing: Third European Conference, ESOC 2014, Manchester, UK, September 2-4, 2014. Proceedings*, pages 171–186. Springer Berlin Heidelberg.
- Caporuscio, M., Inverardi, P., and Pelliccione, P. (2004). Compositional verification of middleware-based software architecture descriptions. In *Proc. of the 26th International Conference on Software Engineering*, ICSE '04, pages 221–230. IEEE Computer Society.
- CloudSoft (2016). brooklyn-tosca. <https://github.com/cloudsoft/brooklyn-tosca>.
- Geraci, A. (1991). *IEEE Standard Computer Dictionary: Compilation of IEEE Standard Computer Glossaries*. IEEE Press.
- Hirmer, P., Breitenbücher, U., Binz, T., and Leymann, F. (2014). Automatic topology completion of TOSCA-based cloud applications. In *INFORMATIK 2014*, volume 232 of *LNI*, pages 247–258. Gesellschaft für Informatik (GI).
- Kopp, O., Binz, T., Breitenbücher, U., and Leymann, F. (2013). Winery – a modeling tool for toasca-based cloud applications. In *Proc. of the 11th International Conference on Service-Oriented Computing*, ICSOC'13, pages 700–704. Springer.
- Marchetti, E. (2014). Foreword of the thematic track: ICT verification and validation. In *Proc. of the 9th International Conference on the Quality of Information and Communications Technology*, QUATIC'14, pages 208–209. IEEE.
- OASIS (2013a). Topology and Orchestration Specification for Cloud Applications. URL: <http://docs.oasis-open.org/tosca/TOSCA/v1.0/TOSCA-v1.0.pdf>.
- OASIS (2013b). Topology and Orchestration Specification for Cloud Applications (TOSCA) Primer. URL: <http://docs.oasis-open.org/tosca/tosca-primer/v1.0/tosca-primer-v1.0.pdf>.
- OASIS (2016). TOSCA Simple Profile in YAML, Version 1.0. URL: <http://docs.oasis-open.org/tosca/TOSCA-Simple-Profile-YAML/v1.0/TOSCA-Simple-Profile-YAML-v1.0.pdf>.
- OpenStack (2016). TOSCA parser. <https://github.com/openstack/tosca-parser>.

- Pelliccione, P., Tivoli, M., Bucchiarone, A., and Polini, A. (2008). An architectural approach to the correct and automatic assembly of evolving component-based systems. *J. Syst. Softw.*, 81(12):2237–2251.
- Speck, A., Pulvermuller, E., Jerger, M., and Franczyk, B. (2002). Component composition validation. *International Journal of Applied Mathematics and Computer Science*, 12(4):581–590.
- Wu, Y., Chen, M.-H., and Offutt, J. (2003). Uml-based integration testing for component-based software. In *Proc. of the 2nd International Conference on COTS-Based Software Systems, ICCBSS'03*, pages 251–260. Springer Berlin Heidelberg.